



gemma4, 16x memory reduction



Qwen3.5, 32x memory reduction



transfers in one token-layer (gemma4)

