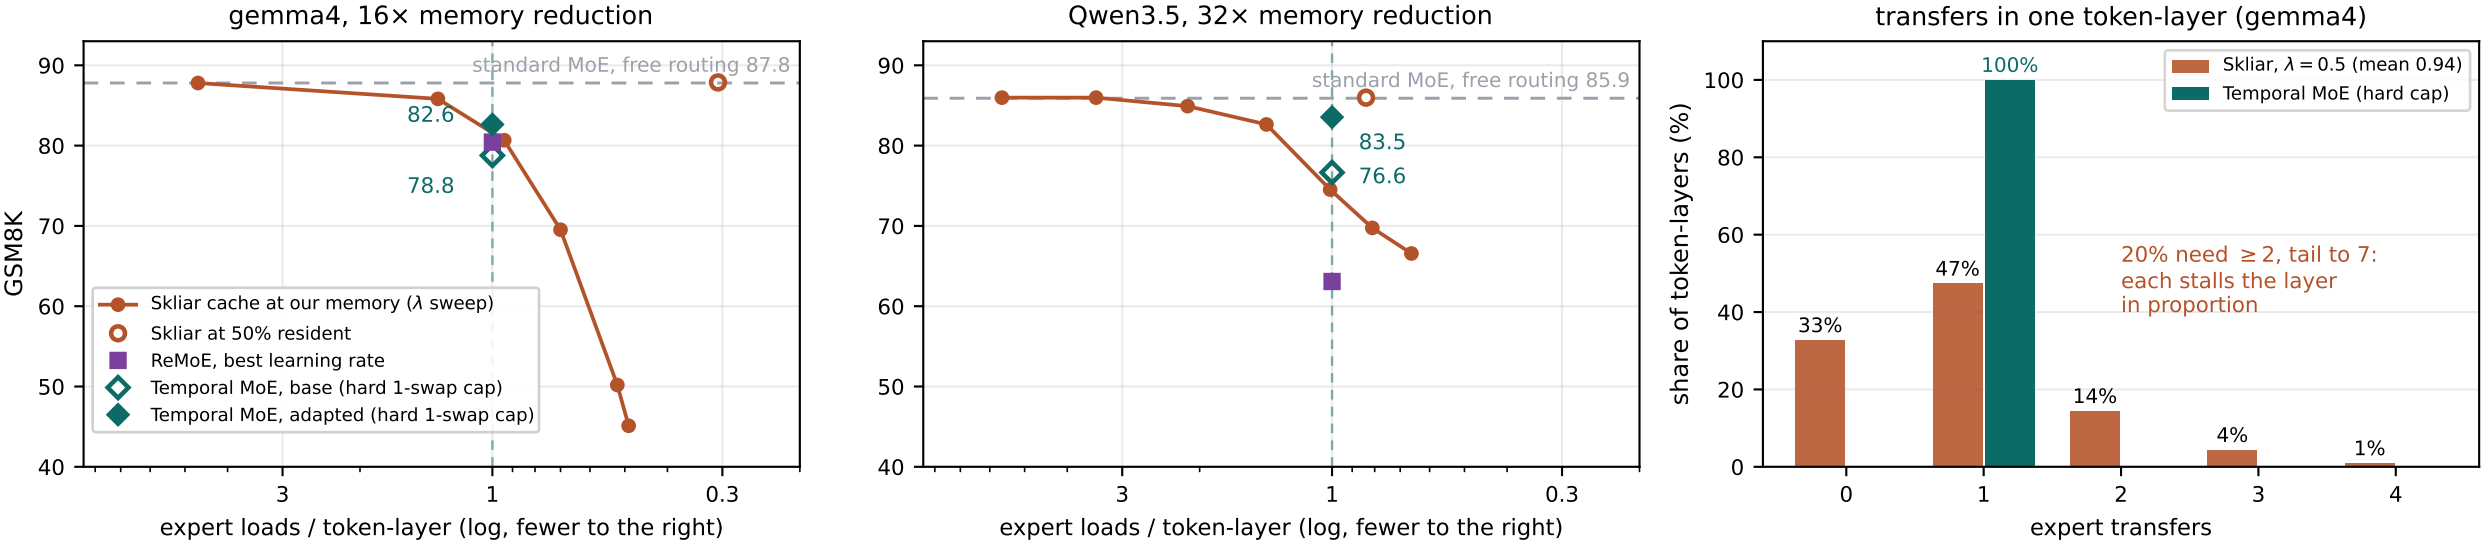
gemma4, 16× memory reduction
standard MoE, free routing 87.8
82.6
78.8
GSM8K
Skliar cache at our memory (λ sweep)
Skliar at 50% resident
ReMoE, best learning rate
Temporal MoE, base (hard 1-swap cap)
Temporal MoE, adapted (hard 1-swap cap)
expert loads / token-layer (log, fewer to the right)
Qwen3.5, 32× memory reduction
standard MoE, free routing 85.9
83.5
76.6
expert loads / token-layer (log, fewer to the right)
transfers in one token-layer (gemma4)
share of token-layers (%)
Skliar, λ = 0.5 (mean 0.94)
Temporal MoE (hard cap)
33%
47%
100%
14%
4%
1%
20% need ≥ 2, tail to 7:
each stalls the layer
in proportion
expert transfers