



Experts active per token (6 of 64)

standard MoE (top-k)



Temporal MoE (resident set used)



unconstrained temporally-trained MoE

